

Cost-Derived Triage of Secret-Scanner Findings When the Credential Cannot Be Verified

Grishma M

grishma.renuka@gmail.com

Abstract

A secret scanner reports that a string in a repository looks like an API key. It does not say whether the key works. The obvious way to find out—calling the provider with the key—means authenticating with a credential belonging to someone else, and we rule it out. This paper studies triage in the *no-direct-verification* setting: the hidden state is not merely unknown, it is one we have decided not to resolve by the available means.

We model the finding as one of three hidden states (live, revoked, fake), price five candidate actions in engineer-minutes, and select by minimising expected cost. No threshold is tuned; the decision boundaries fall out of the cost matrix. Against an idealised escalate-everything baseline, the cost-derived policy saves 23.2% of modelled engineer-minutes per finding while escalating nothing. Every figure here is an expectation under our own cost model, not a measurement of real triage work.

Three results were not designed in. First, the two free features we model cannot distinguish a live key from a revoked one at all: both states carry identical likelihoods on each, so their posterior ratio is invariant at 1.778, which also confines the agent to a one-dimensional slice of the belief simplex. Second, escalation *as a way of resolving the hidden state* is never cost-optimal at any reachable belief: a perfect oracle is worth at most 15.03 minutes there. Third, and least flattering, a policy that observes *nothing at all* already captures 21.2 of the 23.2 percentage points. Joint sensitivity analysis over 40,000 sampled cost models shows the ordering of these results is stable and the magnitudes are not.

1 Introduction

Secret scanners are widely deployed and reasonably good at what they do. What they produce is a finding: *this string, in this file, looks like a credential of this type*. What they do not produce is the thing a responder actually needs, which is whether the credential still works.

The gap matters because the available actions differ enormously in cost. Doing nothing about a live key risks whatever sits behind it. Revoking immediately is cheap and certain, and breaks anything still using it. Rotating safely avoids the outage and costs a deploy cycle. Handing the finding to a person is reliable and spends the scarcest resource in the organisation. The right choice depends on a hidden state, and the standard method for resolving that state—issuing a request with the found credential—is not one we are willing to use against a third party’s account.

That constraint is the origin of the problem we study. It is our choice, not a consensus position: three practitioners told us, unprompted, to simply try the key. We keep it, and treat it as the paper’s most load-bearing assumption.

Contributions.

1. A decision-theoretic formulation of secret-scanner triage in which the decision boundaries are *derived* from a cost matrix rather than tuned. On the slice of the simplex the agent can occupy, the derived boundary sits at $P(\text{live}) = 0.066$ rather than near the 0.5 an engineer would pick by instinct.
2. A structural result: the two free features we model cannot separate live from revoked keys, because both states carry identical likelihoods on each. Given conditional independence this is arithmetic rather than an empirical claim, and it is a property of the features chosen, not of repository evidence in general.
3. A parametric bound on escalation: a perfect oracle is worth at most 15.03 minutes at any reachable belief, so escalation cannot be justified on information-gathering grounds whenever human triage overhead exceeds that. This bounds a human acting as a state oracle; it says nothing about a human who supplies authority or organisational context.
4. A closed-world comparison showing that most of the improvement over the baseline comes from the cost structure rather than from the belief model, and that hand-picked thresholds respond non-monotonically to a prior we cannot pin down.
5. A joint sensitivity study over 40,000 sampled cost and evidence models, identifying which conclusions are robust and which are artefacts of point estimates.

2 Related Work

Detector precision is not a constant. A nine-tool comparison [Basak *et al.*, 2023a] reports precision of about 75% for GitHub’s scanner, 46% for Gitleaks, and below 7% for five of the nine, attributing the spread to over-broad regular expressions and entropy scoring. There is therefore no such thing as *the* base rate of true positives; it is a property of which detector fired. We assume GitHub’s scanner throughout, and note that assuming a weaker detector would change the problem from “what should be done about this key” to “is this junk”.

Credential lifetime. GitGuardian [GitGuardian, 2026] reports that roughly 64% of credentials valid in 2022 were still valid in January 2026. This is a vendor figure and we treat it as a hypothesis, but it points the opposite way to intuition: commit age is weak evidence that a key is dead.

Classification of secrets with language models. Recent work [Islam and others, 2025] reports GPT-4o at 93.29% precision few-shot for deciding whether a string *is* a secret. This is the obvious alternative technology and answers a different question: it improves detection and says nothing about whether a detected credential is live, which is the axis our decision turns on.

Decision theory and the value of information. The formulation here is expected-cost minimisation with pre-posterior analysis, which is classical [Raiffa and Schlaifer, 1961; Howard, 1966]. Deriving a decision threshold from a cost matrix rather than tuning it is the central move of cost-sensitive learning [Elkan, 2001; Domingos, 1999], and the buy-evidence-then-decide structure of our INVESTIGATE action is test-cost-sensitive learning [Turney, 2000]. The sequential formulation we defer to in Section 9 is a POMDP [Smallwood and Sondik, 1973; Kaelbling *et al.*, 1998]. Our contribution is not the machinery but the instance: a domain where the machinery produces predominantly negative results, and where we can say why.

Analyst load. The cost of human attention is the quantity our escalation result turns on. Qualitative work on security operations [Alahmadi *et al.*, 2022] documents analysts discarding the majority of alerts under load, which is evidence against our escalate-everything baseline being what teams actually run.

Revocation is solved for a different credential format. A practitioner observed during our discussion phase that a certificate carries a CRL distribution point—a pointer to its own revocation list—so live-versus-revoked is answered by the credential itself. Public-key infrastructure solved decades ago the exact problem that motivates this paper. An API key has no equivalent. **Our central difficulty is a property of API keys specifically, not of leaked credentials in general.**

3 Problem Formulation

3.1 Input and hidden state

The agent observes one scanner finding: a string in a repository matching a known key format. The hidden state s is one of

- **live** — the credential still authenticates;

Action	What it does
Dismiss	Nothing. The finding is closed.
Investigate	Buy evidence, then decide again.
Escalate	Hand the decision to a human.
Revoke now	Kill the key; accept the outage.
Rotate safely	Replace, migrate, verify, revoke.

Table 1: The five candidate actions. Section 4.5 argues only three are terminal.




Figure 1: The agent. Two free features produce a belief; the cheapest terminal action follows from the cost matrix. Dashed paths are those the policy takes rarely or never: INVESTIGATE fires on 4% of findings, ESCALATE on none, and DISMISS on no belief the agent can reach. That two of five actions go unused is a result, not an omission.

- **revoked** — it authenticated once and no longer does; 138
- **fake** — hard-coded by a developer, never authenticated. 139

We scope to a single provider, OpenAI, and to project keys. This collapses a fourth state that exists elsewhere: Stripe issues distinguishable test keys, OpenAI does not. The trade was deliberate—we gave up a hidden state to gain a probe whose response shape is documented [OpenAI, 2026].

Only the first state carries the exposure risk we price. The other two do not, though “harmless” would be too strong: a revoked key still has audit implications and may have been reused elsewhere.

3.2 Actions

Revoke-now and rotate-safely are kept separate because leaving a system broken in order to be safe is not automatically correct, and a model that cannot express the difference cannot inform the choice. This separation is justified a second time in Section 5.2, from an argument we did not anticipate.

3.3 Baseline and objective

Our baseline is an **idealised escalate-everything policy**: every finding goes to a person, and that person resolves it correctly. We do not claim this is what teams do today—real pipelines allowlist known fixtures, auto-dismiss, and route by pattern before any human sees an alert. It is a deliberately favourable comparator, chosen because it makes the argument clean: **the baseline is never wrong about the hidden state.**

There is therefore no accuracy for an automated method to win, and the only remaining axis is cost.

Feature	live	revoked	fake
Context: pl./ne./pr.	.10/.30/.60	.10/.30/.60	.70/.25/.05
Form: well/mal	.99/.01	.99/.01	.40/.60
last_used: rec./old/null	.60/.20/.20	.02/.78/.20	.01/.04/.95

Table 2: $P(e \mid s)$. Live and revoked are identical on both free features.

Objective. Does a probabilistic, cost-aware policy reach the same decisions as escalate-everything while spending fewer human interruptions—and under what conditions does it stop doing so?

The second clause carries most of the weight. A policy that wins under every assumption we could vary would be evidence the comparison was built badly.

4 Model

4.1 Prior

We *construct* a traceable prior from two external estimates rather than estimating one from representative data. GitHub’s scanner reports about 75% precision; of real credentials, about 64% are still valid. This gives

$$P(\text{live}) = 0.48, \quad P(\text{revoked}) = 0.27, \quad P(\text{fake}) = 0.25.$$

Its virtue is traceability, not accuracy. Two caveats belong here rather than in the limitations. The 0.64 describes credentials known valid in 2022 and still valid in 2026, which is not the population a scanner reports today; and detector precision counts strings that are not credentials at all, whereas our *fake* state is a real-format string that never authenticated. We merge those populations and its likelihood row is a guess for both.

4.2 Evidence

Two features are free—readable from the repository at no cost—and one must be purchased.

Placeholder context (placeholder / neutral / production), from the file path, variable name, and commit message. **Well-formedness** (well-formed / malformed), whether the string matches the published format. **last_used_at** (recent / old / null), a timestamp obtained by asking *our own* account about the key through the provider’s admin API. This never involves authenticating as the suspect key.

Conditional independence. We assume features are conditionally independent given the hidden state, so that $P(c, f, r \mid s) = P(c \mid s)P(f \mid s)P(r \mid s)$. This licenses the multiplication in every posterior below and makes Section 4.3’s result follow from the rows of Table 2 rather than from a joint distribution we never specify. It is an assumption of convenience and plainly imperfect: a string that reads as a placeholder from its path is more likely to be malformed than the product suggests.

4.3 A structural consequence

Note the first two rows of Table 2. **Live and revoked carry identical likelihoods on both free features.** This is not belief, and the agent would probe forever without acting.

	live	revoked	fake
Dismiss	2400	2	2
Investigate	3	3	3
Escalate	142	32	32
Revoke now	332	2	5
Rotate safely	112	30	20

Table 3: The cost matrix, engineer-minutes. The live column is exact component sums.

modelling convenience; it reflects that nothing visible inside a repository distinguishes a key that works from one that was turned off. Identical rows cancel in the update, so the posterior ratio is invariant:

$$\frac{P(\text{live} \mid e_{\text{free}})}{P(\text{revoked} \mid e_{\text{free}})} = \frac{0.48}{0.27} = 1.778 \quad \text{for every } e_{\text{free}}.$$

On our worked case the posterior is $0.6329/0.3560/0.0111$, and $0.6329/0.3560 = 1.778$ to four decimal places. The free evidence does one job excellently—it collapses *fake* from 0.25 to 0.011—and the other not at all.

Neither of these two features can move the ratio that matters, at any value, in any combination. We state it narrowly on purpose: this is a property of the two static text features we chose, not of repository evidence in general. Section 9 names two we did not model—repository visibility and commit semantics—that would plausibly separate the two states, and the claim is safest read as applying to findings in private repositories.

4.4 Costs

Fifteen cells built from nine components, in engineer-minutes, so a practitioner can dispute a component rather than a cell. Components: revoke 2, issue a new key 10, run the probe 3, human attention 30, find the consumers 30 (documented) or 480 (undocumented), deploy 60, verify 10, production outage 240.

Rotate-safely on live is $10 + 30 + 60 + 10 + 2 = 112$; revoke-now is the same hunt and deploy plus the outage, $2 + 240 + 30 + 60 = 332$. Escalate is defined throughout as $30 + \min_a c(a, s)$ —human attention plus whatever that human then correctly does—giving $142/32/32$.

Three cells are direct estimates with no component decomposition: rotating a revoked key (30) or a fake one (20), and revoking a fake one (5). The rotate-on-fake figure is the cost of beginning a rotation and discovering there is nothing to rotate. It is load-bearing: Section 7.5 shows that setting it to zero makes the belief model worth exactly nothing.

4.5 Investigate is not a terminal action

The investigate row is not the same kind of number as the other four. The others are what an action *costs*; this one is an entry fee. Investigate does not close the case—it buys an answer and returns the decision to the agent. Placing 3 into the

minimisation alongside the rest would make it win at every belief, and the agent would probe forever without acting.




Figure 2: Decision regimes on the reachable line, and the six posteriors the agent can actually hold (dots). The derived boundary between revoking and rotating sits at 0.066; the dashed line at 0.5 is where an engineer sets it by instinct. Everything between the two is dismissed by the naive rule and rotated by ours.

We therefore minimise over the four **terminal** actions, and compare investigate separately by whether the expected value of the information exceeds its price. The general rule: *a sequence of steps may be collapsed into a single action when no observation inside the sequence changes what is done next*. Rotate-safely observes but does not branch, so it collapses. Investigate exists only to branch, so it cannot.

5 Policy

5.1 Selection

$$a^* = \arg \min_{a \in \mathcal{A}_{\text{terminal}}} \sum_s P(s | e) c(a, s) \quad (1)$$

No threshold appears in this expression. Thresholds are a consequence of it.

5.2 The derived regimes

With three states, $P(\text{live})$ alone does not determine the cheapest action—the split of the remaining mass between revoked and fake matters too. A boundary is therefore only meaningful relative to a path through the simplex, and Section 4.3 tells us which: the free features lock $P(\text{live}) : P(\text{revoked})$ at 1.778, so the agent walks one line and never leaves it. All boundaries below are measured on that line.

The boundary is at 0.066, not 0.5. An engineer building this by instinct places the line at “more likely than not” and dismisses findings this model rotates; Section 7 measures what that costs.

And **revoke-now owns an entire regime on its own**. We introduced it because outages are not automatically acceptable; the arithmetic justifies it from a different direction, and a model with a single remediation action would never have exposed that region.

5.3 A bound on escalation

At each belief we compute the value of perfect information—the cost saved by an oracle revealing the true state. This

Context	Form	$P(\text{seen})$	$P(\text{live})$	Act	EVSI
placeholder	well	.1442	.3294	rotate	0.00
placeholder	mal	.1057	.0045	revoke	1.40
neutral	well	.2477	.5754	rotate	0.00
neutral	mal	.0398	.0362	revoke	5.21
production	well	.4505	.6329	rotate	0.00
production	mal	.0120	.2400	rotate	0.00

Table 4: Value of the probe at every belief the agent can hold.

upper-bounds what a human who *only resolves the hidden state* could be worth.

Over the whole simplex the maximum is 24.78 minutes, but that maximiser lies at $b \approx (0.12, 0.88, 0.00)$, which the agent cannot reach. Over the six beliefs it can hold,

$$\max_{b \in \mathcal{B}_{\text{reachable}}} \text{VPI}(b) = 15.03 \text{ min.} \quad (2)$$

So escalation is dominated whenever a human costs more than 15.03 minutes. At our assumed 30 it loses comfortably; at the 10–15 minutes a dedicated triage analyst might cost, it does not. **This is a parametric result, not a structural one**, and it bounds only the informational value of a human. A person supplying authority to act on another team’s credential is not bounded by any VPI calculation.

5.4 Uncertainty is not the trigger

It is natural to escalate when confused. We tested the most confused belief available, $(1/3, 1/3, 1/3)$ at 1.585 bits. Rotate-safely costs 54.00; escalate costs 68.67.

Entropy measures confusion. It does not measure whether the confusion is expensive. Two roads at a fork are maximally uncertain and the choice does not matter; a fire alarm that is probably faulty is barely uncertain and urgently needs a person. Entropy retains a role—counting the bits a probe removes—but the escalation trigger is the cost of being wrong.

5.5 When to buy the probe

Three context values and two form values give exactly **six** reachable posteriors, small enough to enumerate rather than sample.

At the four beliefs where rotate-safely is already optimal, every probe outcome leaves it optimal, so the information is worth exactly nothing. Value concentrates at the two beliefs near a boundary.

Because EVSI takes only **two** non-zero values, the probe’s price selects one of three regimes rather than tuning behaviour continuously: below 1.40 it is bought on 14.55% of findings, between 1.40 and 5.21 on 3.98%, and above 5.21 never. We initially priced it at ten minutes by guesswork; a practitioner who uses it offered roughly one minute, and a second warned against concluding anything from a single response. One minute per call, three calls, gives three.

Two things follow. **Investigation becomes rational at 5.21 minutes**, which is a property of the cost matrix rather than of the probe. And **it never becomes important**: even at half a minute the aggregate saving is 0.283 minutes per finding, four tenths of one percent. Information here is sometimes

	Policy	Cost	vs base	$\neq$ P0
base	escalate all	84.80	—	—
P0	prior only	66.86	−21.2%	—
P1	hand-picked, 4 acts	77.79	−8.3%	14.55%
P1-t	hand-picked, 2 acts	181.78	+114.4%	14.55%
P2	cost-derived	65.11	−23.2%	14.55%
P3	P2 + probe	65.03	−23.3%	10.57%

Table 5: Exact expected cost per finding. The last column is the share of findings on which the policy chooses differently from P0’s constant “rotate safely”.

worth buying and never worth much, and those are separate facts.

6 Simulation Study

Six policies on identical findings: the **baseline** (escalate every finding); **P0** (act on the prior, observe nothing); **P1** (free evidence, hand-picked thresholds—dismiss below 0.05, revoke below 0.5, else rotate—over all four terminal actions); **P1-trunc** (the same 0.5 threshold over {rotate, dismiss} only); **P2** (free evidence, minimise expected cost); **P3** (P2 plus the probe when EVSI exceeds its price).

P1 is the comparison that matters: hand-picked thresholds, but the same information and actions available to P2, so it isolates *where the boundaries came from* and nothing else. P1-trunc is an earlier version of that comparison which also restricted the action set; we report it because it produces a dramatic result for the wrong reason.

Expected cost is computed in closed form over all 54 possible worlds and carries no sampling error. Realised cost is measured on 40 cases drawn once with a fixed seed and frozen thereafter.

6.1 Why the study is a simulation, not an experiment

The obvious alternative is SecretBench [Basak *et al.*, 2023b], a labelled corpus of 97,479 candidate secrets, 15,084 confirmed real. We cite it and do not use it, for a specific reason.

SecretBench labels whether a string is a secret. It does not label whether that secret still works. Its fields would ground our *fake* state, the 0.25 of prior mass assigned to it, and both free feature rows. There is no field for live versus revoked—the axis this paper turns on. Fitting to it would sharpen the half of the model that already works and leave the half that matters untouched.

We record the gap as evidence rather than an obstacle. **The largest public corpus of real leaked credentials does not record whether they still authenticate**, and the reason is the one that motivates this paper: establishing it would mean testing each key against its provider. The label is missing from the dataset for precisely the reason the state is hidden from our agent.

7 Results

On the 40 frozen cases: baseline 2820 minutes, P2 2018, 28.4% saving. That draw contains 14 live keys out of 40,

live share of 0.35 against a prior of 0.48, so the sampled and exact figures are not interchangeable.

7.1 The cost structure does the work

P0 observes nothing whatsoever. It knows the base rates, concludes rotate-safely is cheapest in expectation, and applies it to every finding. It beats the baseline by 21.2%. The full belief model reaches 23.2%. **The evidence is worth 1.75 minutes per finding, and the remaining twenty-one points come from a single decision: stop asking a human.** It does change the action on 14.55% of findings—it is not inert—but the action it changes to is cheap enough that the aggregate barely moves.

The mechanism is that rotate-safely costs 112/30/20: it is nearly state-blind and tolerable whichever state obtains. When a hedge that good is available, identifying the state buys very little. This is the EVSI result one level up.

7.2 What hand-picked thresholds actually cost

An earlier version of this comparison reported that the naive 0.5 threshold costs more than twice the baseline, and treated that as the strongest argument for deriving boundaries. **That comparison was confounded.** P1-trunc changes two things at once: it moves the boundary to 0.5 *and* restricts the action set, withholding revoke-now. Holding the threshold fixed and varying only the fallback: dismiss 181.78 (+114.4%), escalate 71.13 (−16.1%), revoke 74.25 (−12.4%), rotate 66.86 (−21.2%).

The catastrophe is caused by dismissing, not by thresholding. Given four actions and a second cut at 0.05, P1 costs 77.79—8.3% better than the baseline and 19.5% worse than P2.

What survives is a difference in response to a prior we cannot pin down. Across live shares 0.40 to 0.80, P1 runs 103.8, 129.2, 77.8, 83.1, 91.9 while P2 runs 50.2, 56.4, 65.1, 68.8, 75.1. P1 is **non-monotonic**: it gets more expensive as the world gets safer, and at a live share of 0.40 costs 60% *more* than the baseline it beat at 0.64. **A fixed boundary cannot track a moving prior, whereas one read off the cost matrix moves when the costs move.**

7.3 Evidence is worthless when one action wins everywhere

Priced at zero outage, revoke-now is optimal at every reachable belief, and P0 and P2 then cost *exactly* 46.0. Reading the features changes nothing, because nothing they could say would change the action.

7.4 Information concentrates at decision boundaries

The probe is worth 1.0 minutes when an outage costs 60 and 0.08 when it costs 240—twelve times more valuable when a cheap outage puts revoke-now and rotate-safely into close competition.

The converse refutes an extension we tested. Probing *which systems consume the credential* has no value for choosing between the remediation actions: the hunt appears identical in both, $c(\text{rotate}, \text{live}) = 82 + F$ and $c(\text{revoke}, \text{live}) =$

Chose	On a key that was	Cases	Regret
rotate safely	revoked	11	308
rotate safely	fake	3	54
revoke now	fake	12	36
(correct)	—	14	0

Table 6: Every error is over-remediation. Not one case in forty under-reacts.

302 + F , so their difference is 220 for every F and the boundary sits at 0.06588 throughout. **Information that changes what you pay, but not the ordering of what you might do, is worth nothing.**

The same invariance disposes of a related worry: scaling the cost of dismissing a live key from 2,400 to 120,000 leaves the revoke/rotate boundary at 0.06588, because that component appears in neither action.

7.5 The belief model’s value lives in one cost cell

Varying the cost of rotating a *fake* key: at 0, P0 and P2 are identical at 61.86 and the belief model is worth nothing at all; at 20 the gap is 1.75; at 120 it is 14.18. Every result in the preceding subsection therefore rests on a single estimated cell.

7.6 Where the policy loses, and what it never does

Defining an incorrect decision as one a hindsight-perfect agent would have made differently, with regret $c(a, s) - \min_{a'} c(a', s)$, P2 pays 2018 minutes against a hindsight-optimal 1620: total regret 398 minutes, 19.7%.

The five costliest errors are the *same* error five times: a well-formed key in a real-looking file, rotated, already revoked, 28 minutes wasted each. The agent has no idiosyncratic failures, only one systematic bias—the invariant ratio arriving as a bill. Those decisions were correct given the information and wrong given the world.

The two expensive errors available—dismissing a live key (regret 2288) and revoking one (220)—did not occur. The first cannot: P2 never dismisses at a reachable belief. The second can, and roughly two were expected on this draw; zero occurred, which is a sampling accident rather than a property.

7.7 Joint sensitivity

Every sweep above varies one component and holds twelve fixed, which flatters the model. We sampled the whole specification jointly—40,000 draws, each quantity from a two-piece lognormal whose most-likely value is the median and whose stated low and high are the 5th and 95th percentiles—and ran it twice: once with the evidence model fixed, once with every likelihood row also perturbed (Dirichlet, $\alpha = 40$), the live and revoked rows independently, deliberately breaking the property behind Section 4.3.

Perturbing the evidence model breaks the invariance and changes almost nothing. The live-to-revoked log-ratio drift across buckets rises from exactly 0.000 to a median of 0.88—lower than an earlier draft made. Deriving a boundary 2.53—a factor of roughly twelve, meaning the free features rather than choosing it does not make the policy dramatically do separate the two states in most of those worlds—and no cheaper—a hand-picked threshold with a sensible fallback

Claim	costs	+ evidence
P2 costs less than the baseline	98.02%	98.64%
Boundary below 0.5	93.26%	93.24%
Belief model worth < 5 min	91.34%	88.13%
P0 also beats the baseline	79.20%	79.20%
Escalate chosen nowhere	70.29%	69.75%
Probe bought somewhere	59.29%	57.98%

Table 7: Share of sampled draws in which each claim holds. These are not confidence levels, and components are drawn independently.




Figure 3: Escalation is a function of one number we guessed. With every other cost fixed at its point estimate, the agent escalates every finding when a human costs three minutes and none when a human costs twenty. Our assumed thirty is well past the crossover; a dedicated triage analyst would not be.

conclusion moves by more than 3.2 percentage points. Section 4.3’s invariance is an idealisation, and the conclusions that appear to rest on it do not need it.

Escalation, tested rather than observed. Figure 3 holds every other component at its point estimate and varies only human attention. Escalation goes from every finding to none across a twenty-minute range. **Section 5.3’s result is a consequence of pricing human attention at 30 minutes**, not of the structure of the problem, and we withdraw any reading of it as structural.

Every point estimate sits near its median, so our cost model is not an outlier among plausible ones; the intervals are nonetheless very wide. **The claims worth keeping are about ordering rather than magnitude.**

8 Discussion

Three of this paper’s results are negative, and we think that is the point. The free features cannot separate the states that matter; escalation is not cost-optimal at our assumed human cost; the probe is usually not worth buying. Each was established by computation rather than assertion, and each was contrary to our expectation before it was computed.

The methodological observation we would defend is narrower than an earlier draft made. Deriving a boundary rather than choosing it does not make the policy dramatically cheaper—a hand-picked threshold with a sensible fallback




Figure 4: Point estimates against their sampled distributions, 5th to 95th percentile. Every value we quote sits near the median of its distribution, so the cost model is not an outlier among plausible ones—but the intervals are wide enough that quoting any single figure implies a precision the model does not have. A negative boundary means revoke-now never wins in that world.

beats the baseline too. What it changes is **what an objection looks like**. A tuned threshold can only be defended by authority; a derived one can only be attacked through a cost, which is a claim a practitioner can correct. One did, by an order of magnitude. And a derived boundary moves when the world moves, which is why P2 responds monotonically to the prior where P1 does not.

The result we would most like to see contradicted is Section 7’s: if the belief model is worth 1.75 minutes, most of what makes this agent useful is the decision to stop escalating, which requires no probabilistic machinery. Section 7.5 shows that rests on a single estimated cell, and at zero that cell makes the belief model worth precisely nothing. The finding and its weakest input are the same input.

9 Limitations

The refusal to test the key is ours alone. Three practitioners advised calling the provider with the found credential. Every downstream result exists because of a choice the field does not share.

Eight of nine cost components are our estimates. The ninth was corrected by practitioners and changed a result, which is the best available evidence that the other eight are also wrong in unknown directions.

The case generator draws from the agent’s own likelihood tables, so the agent is evaluated in a world that agrees with its assumptions. This flatters every policy that uses evidence.

Actions are atomic; the world is sequential. Rotate safely is really six steps with evidence arriving partway through, and a sequential formulation would find value for investigation this model cannot see.

Rotate-safely is modelled as deterministic remediation. With probability ϵ a rotation silently fails on an undocumented consumer and the outage is paid anyway; introducing ϵ makes the revoke/rotate boundary depend on the outage cost and destroys the invariance of Section 7. A practitioner described exactly this failure.

Escalation may be a permission boundary rather than a priced option. An agent may lack the authority to revoke another team’s credential, which is a constraint on the action space and not a candidate inside it.

The free feature set is not the whole repository. Repository visibility and commit semantics would plausibly separate live from revoked; Section 7.7 shows that even when they do, the conclusions barely move.

The probe requires a privileged credential. Reading `last_used.at` means the scanning pipeline holds an organisation-wide admin key to triage a single project key—a high-severity credential introduced to resolve a low-severity finding, and the model prices the minutes without pricing that asymmetry.

A fourth state is known and unmodelled. A credential that authenticates but is authorised for nothing—a 403 rather than a 401—is live by our definition yet carries none of the risk that makes dismissing a live key cost 2400.

No live call was made. No admin credential was held, and the repository contains no credential material; every string is a synthetic feature descriptor.

10 Conclusion

We formulated secret-scanner triage as a decision under uncertainty in which the hidden state cannot be observed without doing something we are unwilling to do, and showed the decision is nonetheless tractable because the *costs* are knowable even when the state is not.

The policy saves 23.2% of modelled engineer-minutes while never escalating. But the honest headline is smaller and more interesting: a policy observing nothing at all captures 21.2 of those points, and joint sensitivity analysis puts the saving anywhere between 0.5% and 51%. In this domain a safe action that is cheap in every state does most of the work a belief model appears to be doing—and the contribution of decision theory here is less the posterior than the discipline of pricing the actions before choosing between them.

Acknowledgements

Eight practitioners on `r/sysadmin` and `r/devops` answered questions that changed this work: one repriced the probe by an order of magnitude, one withdrew an assumption about rotation verifying itself, and one established that live-versus-revoked is solved for certificates and unsolved for API keys.

References

[Alahmadi *et al.*, 2022] Bushra A. Alahmadi, Louise Axon, and Ivan Martinovic. 99% false positives: A qualitative study of SOC analysts’ perspectives on security alarms. In *Proceedings of the 31st USENIX Security Symposium*, 2022.

577 [Basak *et al.*, 2023a] Setu Kumar Basak, Lorenzo Neil,
578 Bradley Reaves, and Laurie Williams. A comparative
579 study of software secrets reporting by secret detection
580 tools. *arXiv preprint arXiv:2307.00714*, 2023. Verify au-
581 thor list against the paper before submission.

582 [Basak *et al.*, 2023b] Setu Kumar Basak, Lorenzo Neil,
583 Bradley Reaves, and Laurie Williams. SecretBench: A
584 dataset of software secrets. In *Proceedings of the 20th In-*
585 *ternational Conference on Mining Software Repositories*
586 *(MSR)*, 2023. arXiv:2303.06729. Verify author list before
587 submission.

588 [Domingos, 1999] Pedro Domingos. MetaCost: A general
589 method for making classifiers cost-sensitive. In *Proceed-*
590 *ings of the 5th ACM SIGKDD International Conference on*
591 *Knowledge Discovery and Data Mining*, pages 155–164,
592 1999.

593 [Elkan, 2001] Charles Elkan. The foundations of cost-
594 sensitive learning. In *Proceedings of the 17th International*
595 *Joint Conference on Artificial Intelligence (IJCAI)*, pages
596 973–978, 2001.

597 [GitGuardian, 2026] GitGuardian. The state of secrets
598 sprawl 2026. Technical report, GitGuardian, 2026. Vendor
599 report.

600 [Howard, 1966] Ronald A. Howard. Information value the-
601 ory. *IEEE Transactions on Systems Science and Cybernet-*
602 *ics*, 2(1):22–26, 1966.

603 [Islam and others, 2025] Md. Nazmul Islam et al. Secret
604 breach detection in source code with large language mod-
605 els. *arXiv preprint arXiv:2504.18784*, 2025. Author list
606 needs checking against the paper.

607 [Kaelbling *et al.*, 1998] Leslie Pack Kaelbling, Michael L.
608 Littman, and Anthony R. Cassandra. Planning and act-
609 ing in partially observable stochastic domains. *Artificial*
610 *Intelligence*, 101(1–2):99–134, 1998.

611 [OpenAI, 2026] OpenAI. Administration API — project
612 API keys. OpenAI Documentation, 2026.

613 [Raiffa and Schlaifer, 1961] Howard Raiffa and Robert
614 Schlaifer. *Applied Statistical Decision Theory*. Harvard
615 University Graduate School of Business Administration,
616 1961.

617 [Smallwood and Sondik, 1973] Richard D. Smallwood and
618 Edward J. Sondik. The optimal control of partially observ-
619 able markov processes over a finite horizon. *Operations*
620 *Research*, 21(5):1071–1088, 1973.

[Turney, 2000] Peter D. Turney. Types of cost in inductive
621 concept learning. *Proceedings of the Workshop on Cost-*
622 *Sensitive Learning, ICML*, 2000.
623